

## 1. Output of cProfile, showing the top 10 most time-consuming functions:

```
(base) linnekst@linnekst:/mnt/c/Users/linne/advanced_programming/day2-bestpractices-1$ python -m pstats matmult.prof
Welcome to the profile statistics browser.
matmult.prof% sort tottime
matmult.prof% stats 10
Tue Mar  4 15:31:08 2025      matmult.prof

      1101421 function calls (1101394 primitive calls) in 2.657 seconds

Ordered by: internal time
List reduced from 118 to 10 due to restriction <10>

ncalls  tottime  percall  cumtime  percall filename:lineno(function)
   1      2.476    2.476    2.657    2.657 matmult.py:1(<module>)
125250   0.061    0.000    0.119    0.000 /home/linnekst/mambaforge/lib/python3.10/random.py:292(randrange)
125250   0.035    0.000    0.047    0.000 /home/linnekst/mambaforge/lib/python3.10/random.py:239(_randbelow_with_getrandbits)
125250   0.020    0.000    0.139    0.000 /home/linnekst/mambaforge/lib/python3.10/random.py:366(randint)
   250    0.015    0.000    0.015    0.000 {built-in method builtins.print}
375750   0.012    0.000    0.012    0.000 {built-in method _operator.index}
   250    0.009    0.000    0.079    0.000 matmult.py:9(<listcomp>)
   250    0.009    0.000    0.078    0.000 matmult.py:14(<listcomp>)
158751   0.007    0.000    0.007    0.000 {method 'getrandbits' of '_random.Random' objects}
125250   0.005    0.000    0.005    0.000 {method 'bit_length' of 'int' objects}

matmult.prof% |
```

- Line 1 in matmult.py takes the most time (2.476s total) and should hence be where we start the speed optimization. matmult.py runs all executions at the top level (there is no function defined in the script) which is the cause for the cProfile output attributing the time to Line 1.

Other observations:

- random.randint(), randrange(), and \_randbelow\_with\_getrandbits() each appear 125,250 times, indicating the random number generation also is time consuming.
- List comprehension (line 9 and 14) also takes some time

To analyze the matmult.py with line\_profiler we need an actual function in the script to be analyzed. The output shows the time allocated to each of the lines in this function. The results show that the most time-consuming line is line 19, which is the matrix multiplication.

```
Total time: 7.57198 s
File: matmult_modified.py
Function: matrix_multiplication at line 14

Line #      Hits          Time    Per Hit    % Time  Line Contents
=====
   14              1              0.0      0.0      0.0      @profile
   15              1              0.0      0.0      0.0      def matrix_multiplication():
   16          251          37.9      0.2      0.0          for i in range(len(X)): # Iterate through rows of X
   17        63000         6990.3      0.1      0.1              for j in range(len(Y[0])): # Iterate through columns of Y
   18      15750250      1772190.3      0.1     23.4                  for k in range(len(Y)): # Iterate through rows of Y
   19      15687500      5792762.8      0.4     76.5                      result[i][j] += X[i][k] * Y[k][j]
```

## 2. Investigate the performance of the eulers72.py script. In which line(s) of the script would you start optimizing for speed?

Output from cProfile:

```
(base) linnekst@linnekst:/mnt/c/Users/linne/advanced_programming/day2-bestpractices-1$ python -m pstats euler72.prof
Welcome to the profile statistics browser.
euler72.prof% sort tottime
euler72.prof% stats 10
Tue Mar  4 16:31:42 2025      euler72.prof

      161691 function calls in 0.066 seconds

Ordered by: internal time
List reduced from 77 to 10 due to restriction <10>

ncalls  tottime  percall  cumtime  percall  filename:lineno(function)
    9999    0.034    0.000    0.044    0.000  euler72.py:22(factorize)
    9999    0.013    0.000    0.058    0.000  euler72.py:51(fast_phi)
   99314    0.008    0.000    0.008    0.000  {built-in method math.sqrt}
      1    0.004    0.004    0.066    0.066  euler72.py:1(<module>)
   32153    0.003    0.000    0.003    0.000  {method 'append' of 'list' objects}
      1    0.002    0.002    0.002    0.002  euler72.py:5(gen_primes)
      5    0.001    0.000    0.001    0.000  {built-in method posix.stat}
   10000    0.001    0.000    0.001    0.000  {built-in method builtins.len}
      1    0.000    0.000    0.000    0.000  {built-in method _imp.create_dynamic}
      4    0.000    0.000    0.002    0.000  <frozen importlib._bootstrap_external>:1536(find_spec)
```

- The factorization-function at line 22 and the fast\_phi function at line 51 are the most time-consuming functions and where I should start optimizing the script for speed.

Output from line\_profiler with one @profile decorator at each function

```
(base) linnekst@linnekst: /mnt/c/Users/linne/advanced_programming/day2-bestpractices-1$ kernprof -l -v euler72.py
30397485.0
Wrote profile results to euler72.py.lprof
Timer unit: 1e-06 s
```

```
Total time: 0.0061851 s
File: euler72.py
Function: gen_primes at line 4
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
4					@profile
5					def gen_primes(n):
6	1	5.4	5.4	0.1	l = range(2,n)
7	1	0.5	0.5	0.0	primes = []
8	999	349.4	0.3	5.6	for j in range(0,len(l)):
9	998	249.8	0.3	4.0	p = True
10	2968	851.7	0.3	13.8	for d in primes:
11	2967	1978.0	0.7	32.0	if(d > sqrt(l[j])):
12	167	53.3	0.3	0.9	break
13	2800	1720.5	0.6	27.8	if(l[j] % d == 0):
14	830	201.1	0.2	3.3	p = False
15	830	222.2	0.3	3.6	break;
16	998	413.2	0.4	6.7	if(p):
17	168	139.3	0.8	2.3	primes.append(l[j])
18					
19	1	0.5	0.5	0.0	return primes

```
Total time: 0.0843567 s
File: euler72.py
Function: factorize at line 21
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
21					@profile
22					def factorize(n,primes):
23	9999	1368.0	0.1	1.6	factors = []
24	9999	1286.4	0.1	1.5	init_n = n
25	96347	13360.3	0.1	15.8	for p in primes:
26	118736	27531.0	0.2	32.6	while(n%p == 0):
27	22389	3923.1	0.2	4.7	n = n/p
28	22389	5269.9	0.2	6.2	factors.append(p)
29	96347	24758.1	0.3	29.3	if(p > sqrt(n)):
30	9999	1458.7	0.1	1.7	break
31	9999	1718.5	0.2	2.0	if(n > 1):
32	9596	2236.0	0.2	2.7	factors.append(n)
33	9999	1446.8	0.1	1.7	return factors

```
Total time: 0.194258 s
File: euler72.py
Function: fast_phi at line 50
```

Line #	Hits	Time	Per Hit	% Time	Line Contents
50					@profile
51					def fast_phi(n,primes):
52	9999	169224.5	16.9	87.1	factors = factorize(n,primes)
53	9999	2223.0	0.2	1.1	phi = factors[0]-1
54	31985	7963.3	0.2	4.1	for i in range(1,len(factors)):
55	21986	5908.1	0.3	3.0	if(factors[i] == factors[i-1]):
56	7685	3574.9	0.5	1.8	phi *= (factors[i]-1)*(factors[i])/(factors[i]-1)
57					else:
58	14301	4122.8	0.3	2.1	phi *= (factors[i]-1)
59	9999	1241.5	0.1	0.6	return phi

- The most time consuming row in the factorize function is row 26, while(n%p == 0) was executed 118736 times.
- The most time consuming row in the fast\_phi function is row 52, factors = factorize(n,primes), which used 87% of the time spent by this function.

3. Using NumPy for the matrix multiplication reduces the running time.

```
(base) linnekst@linnekst:/mnt/c/Users/linne/advanced_programming/day2-bestpractices-1$ python -m pstats npmatmult.prof
Welcome to the profile statistics browser.
npmatmult.prof% sort tottime
npmatmult.prof% stats 10
Tue Mar  4 16:54:53 2025      npmatmult.prof
```

99120 function calls (96458 primitive calls) in 0.187 seconds

Ordered by: internal time  
List reduced from 972 to 10 due to restriction <10>

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
645	0.040	0.000	0.040	0.000	{built-in method posix.stat}
28/27	0.039	0.001	0.056	0.002	{built-in method _imp.create_dynamic}
1	0.012	0.012	0.012	0.012	npmatmult.py:13(matrix_multiplication)
117	0.010	0.000	0.010	0.000	{built-in method marshal.loads}
28/21	0.004	0.000	0.016	0.001	{built-in method _imp.exec_dynamic}
313	0.003	0.000	0.049	0.000	<frozen importlib._bootstrap_external>:1536(find_spec)
239/238	0.003	0.000	0.006	0.000	{built-in method builtins.__build_class__}
117	0.002	0.000	0.002	0.000	{method 'read' of '_io.BufferedReader' objects}
11958	0.002	0.000	0.002	0.000	{built-in method builtins.isinstance}
1401	0.002	0.000	0.003	0.000	/home/linnekst/mambaforge/lib/python3.10/typing.py:986(__setattr__)

npmatmult.prof% quit

Goodbye.

(base) linnekst@linnekst:/mnt/c/Users/linne/advanced\_programming/day2-bestpractices-1\$ nano npmatmult.py

(base) linnekst@linnekst:/mnt/c/Users/linne/advanced\_programming/day2-bestpractices-1\$ kernprof -l -v npmatmult.py

Wrote profile results to npmatmult.py.lprof

Timer unit: 1e-06 s

Total time: 0.0123502 s

File: npmatmult.py

Function: matrix\_multiplication at line 13

Line #	Hits	Time	Per Hit	% Time	Line Contents
13					@profile
14					def matrix_multiplication():
15	1	12349.8	12349.8	100.0	result = np.dot(X, Y)
16	1	0.4	0.4	0.0	return result